

Generated text is not authoritative evidence

Frozen research artefact — no live model call

This talk evaluates a governed workflow that keeps machine candidates, automated checks, human validation and source evidence visibly separate — not a deployed decision engine.

Machine candidate

Automated check

Human review

Source evidence

EEEM004 viva presentation · UK COVID-19 Inquiry Module 2 case study
Not live, not deployed, not automatically authoritative.

Main question: Can an LLM-assisted workflow produce useful, auditable candidates without treating generated text as authoritative evidence?

Inquiry transcripts hide decisions in long, mixed discourse

Public inquiry text holds decision evidence that is lengthy, dispersed and easy to mis-extract as if authoritative.

Question

Procedure

Measure

Recall

Agreements sit beside questions and procedure.

“Decision” language is mixed and dispersed.

Naïve LLM summaries can look authoritative.

Aim: traceable candidates under human oversight

Exact dissertation aim

To design and evaluate an LLM-assisted method for creating traceable candidate decision-journal entries from complex public records, using UK COVID-19 Inquiry Module 2 transcripts as a case study.

Presentation question: Can an LLM-assisted workflow produce useful, auditable decision-journal candidates from public inquiry transcripts without treating generated text as authoritative evidence?

1 Develop source-linked extraction

2 Evaluate on eight hearings

3 Freeze 414-candidate journal

4 Validate n=50 validity/strength


5 Organise + n=60 interpret/faithfulness

6 Examine limits and scaling

No numbered research questions in Chapter 1 — Aim + Objectives 1–6 only.


Generation, checks, human review and source stay separate

Later stages do not retrospectively turn machine-generated text into authoritative evidence.



Machine generation does not become authoritative because later checks exist.

Research prototype — not production deployment.



Traceable candidates are common; journal-ready ones are not

Layered evaluation — no single F1. Pass rates $\neq$ acceptance. no $\times$ high and moderate κ matter more than volume.

Scale & traceability

414

candidates

351/414

traceability pass

1/6 vs 5/6

keyword vs LLM

Validity $\neq$ evidence strength

21/50

no $\times$ high

0.48

rule κ vs B

0.39

LLM κ vs B

Framework & faithfulness

11/60

JEE mapped

37/60

DQ mapped

8/25/20/7

faithfulness

Candidates are not accepted decisions. · Supporting: 5/10/0 triangulation · 20 clusters · 50/53 · 49/50

Quote support ≠ journal membership (phase1-082)

SOURCE EVIDENCE

The hearing adjourned until 10 am on Friday, 1 December 2023

FROZEN LLM CANDIDATE

The hearing will resume at 10 o'clock tomorrow.

Traceability: PASS

Rubric B = HIGH (supported wording)

Rubric A = NO (not journal-valid)

Flag: procedural hearing administration · ID phase1-082

**Quotation support does not by itself establish decision-journal validity.
Supported wording ≠ valid journal entry.**

Contribution is governance of the AI–source–human route

Not a new SOTA extraction algorithm — a governed review process.

Methodological

Link candidates to sources; separate validity, strength and faithfulness.

Artefact

Fixed 414-entry journal plus offline examiner demonstration.

Empirical

Bounded $n=414$ / $n=50$ / $n=60$ feasibility evidence.

Governance

Keep machine, automated, human and source layers inspectable.

Make the route from source to interpretation inspectable

Limits carried forward: One inquiry · one model · single reviewer · Frozen outputs · moderate κ · interpretive mappings

The system does not replace authoritative inquiry evidence or human policy judgement. It makes the route from source evidence to machine-generated candidate and human interpretation inspectable and auditable.

Questions

Presentation complete without live demo · optional offline walkthrough ≤ 2 min after the centrepiece only if time remains.